

INTENT

Prepare the dataset for analysis, using methods learnt and researched independently, including but not limited to Missing Value Imputation, Outlier Detection etc. Find out your observations on different variables using descriptive statistics, Visualization etc. Report if there is any pattern present in the data. Take LC50 [-LOG(mol/L)] as the target variable and fit different models for regression using other variables present in the data. Optimize the model parameters and find the best performing model. Compare all models used, using various performance metrics.

Provide inferences to your findings. Data Set Information: LC50 data, which is the concentration that causes death in 50% of test fish over a test duration of 96 hours, was used as a model response. The data comprised 6 molecular descriptors: MLOGP (molecular properties), CIC0 (information indices), GATS1i (2D autocorrelations), NdssC (atom-type counts), NdsCH ((atom-type counts), SM1_Dz(Z) (2D matrix-based descriptors). Details can be found in the quoted reference: M. Cassotti, D. Ballabio, R. Todeschini, V. Consonni. Attribute Information: 6 molecular descriptors and 1 quantitative experimental response: 1) CIC0 2) SM1_Dz(Z) : 0 means missing value 3) GATS1i 4) NdsCH 5) NdssC 6) MLOGP 7) quantitative response, LC50 [-LOG(mol/L)]

```
In [59]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

data =pd.read_csv("qsar_fish_toxicity.csv")

data.info()
data.head(10)
```

dtypes: float64(6), int64(1)
memory usage: 49.8 KB

```
Out[59]:
```

	CIC0	SM1_Dz(Z)	GATS1i	NdsCH	NdssC	MLOGP	LC50 [-LOG(mol/L)]
0	3.260	0.829	1.676	0.0	1	1.453	3.770
1	2.189	0.580	0.863	0.0	0	1.348	3.115
2	2.125	0.638	0.831	0.0	0	1.348	3.531
3	3.027	0.331	1.472	1.0	0	1.807	3.510
4	2.094	0.827	0.860	0.0	0	1.886	5.390
5	3.222	0.331	2.177	0.0	0	0.706	1.819
6	3.179	0.000	1.063	0.0	0	2.942	3.947
7	3.000	0.000	0.938	1.0	0	2.851	3.513
8	2.620	0.499	0.990	0.0	0	2.942	4.402
9	2.834	0.134	0.950	0.0	0	1.591	3.021

Remove Null values

```
In [40]: columns = ['SM1_Dz(Z)']
data[columns] = data[columns].replace(0, float('NaN'))
data.info()
data.isna().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 908 entries, 0 to 907
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CIC0            906 non-null    float64
 1   SM1_Dz(Z)       870 non-null    float64
 2   GATS1i          906 non-null    float64
 3   NdsCH           907 non-null    float64
 4   NdssC           908 non-null    int64  
 5   MLOGP           905 non-null    float64
 6   LC50 [-LOG(mol/L)] 906 non-null  float64
dtypes: float64(6), int64(1)
memory usage: 49.8 KB
```

```
Out[40]: CIC0                2
          SM1_Dz(Z)        38
          GATS1i           1
          NdsCH             0
          NdssC            0
          MLOGP             3
          [-LOG(mol/L)]    2
dtype: int64
```

Calculating null values

```
In [41]: : data.isna().sum()/len(data)*100
          data.skew()
```

```
Out[41]: : CIC0            0.045111
          SM1_Dz(Z)        0.736832
          GATS1i          3.392500
          NdsCH           7.489497
          NdssC            -0.038095
          MLOGP           0.254314
          [-LOG(mol/L)]    0.254314
          dtype: float64
```

After obtaining unsatisfactory results from the previous method, we made the decision to utilize the KNN Imputer for handling null values. Upon observing the model's performance, we concluded that this approach is more suitable for replacing missing data. The KNN Imputer

leverages neighboring data points to impute missing values, making it a robust technique that

KNN method

Based on unsatisfactory results from the previous method, we decided to use the KNNImputer to handle null values. Observing the model's performance, we concluded that this approach is more suitable for replacing missing data. The KNN Imputer uses neighboring data points to impute missing values, making it a robust technique that improves accuracy and reliability of our analysis. This choice is expected to enhance our models and lead to meaningful insights and predictions.

In [42]:

```
# Apply K means to missing value treatment
from sklearn.impute import KNNImputer
imputer = KNNImputer(n_neighbors=5)
df_imputed = pd.DataFrame(imputer.fit_transform(data), columns=data.columns)
df_imputed['NdssC'] = df_imputed['NdssC'].astype(int)
print(df_imputed.isnull().sum())
data = df_imputed
data.isna().sum()/len(data)*100
```

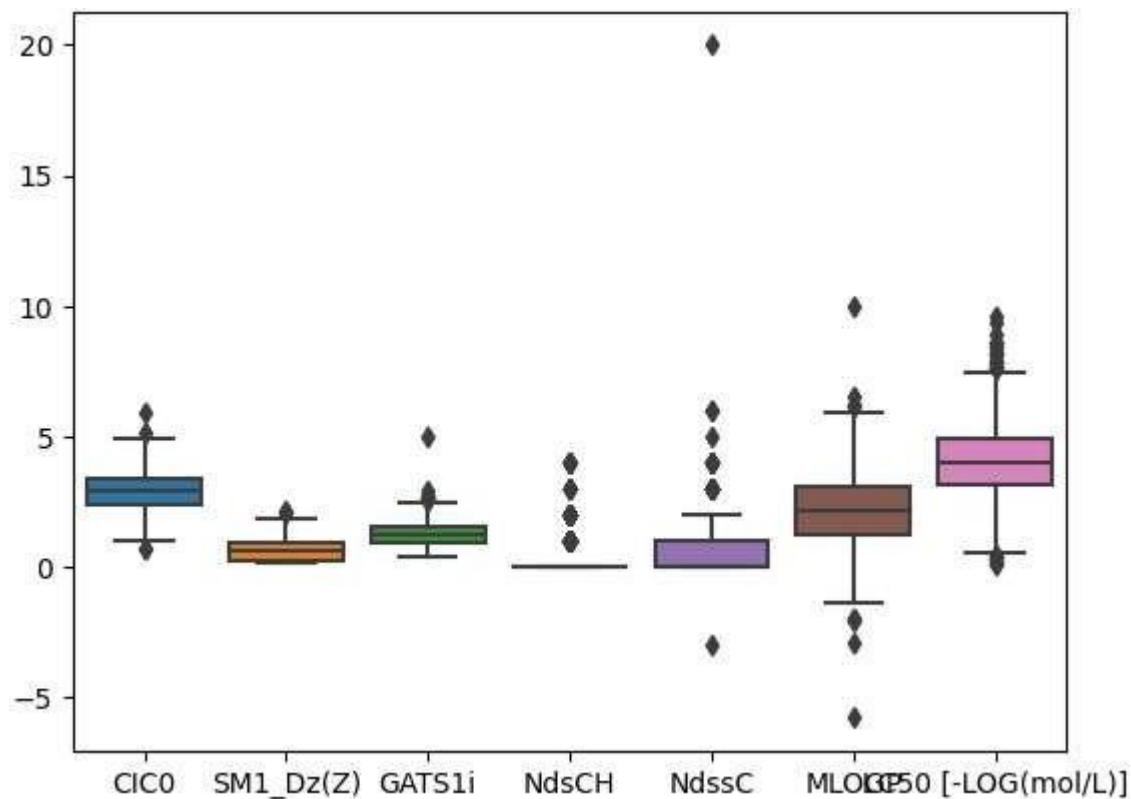
```
CICO          0
SM1_Dz(Z)      0
GATS1i         0
NdssCH         0
NdssC          0
MLOGP          0
[-LOG(mol/L)]  0
dtype: int64
```

Out[42]

```
: CICO          0.0
SM1_Dz(Z)      0.0
GATS1i         0.0
NdssC          0.0
MLOGP          0.0
[-LOG(mol/L)]  0.0
dtype: float64
```

```
In [43]: sns.boxplot(data=data[["CIC0", "SM1_Dz(Z)", "GATS1i", "NdsCH", "NdssC", "MLOGP"]])
```

```
Out[43]: <Axes: >
```



Removing Outliers

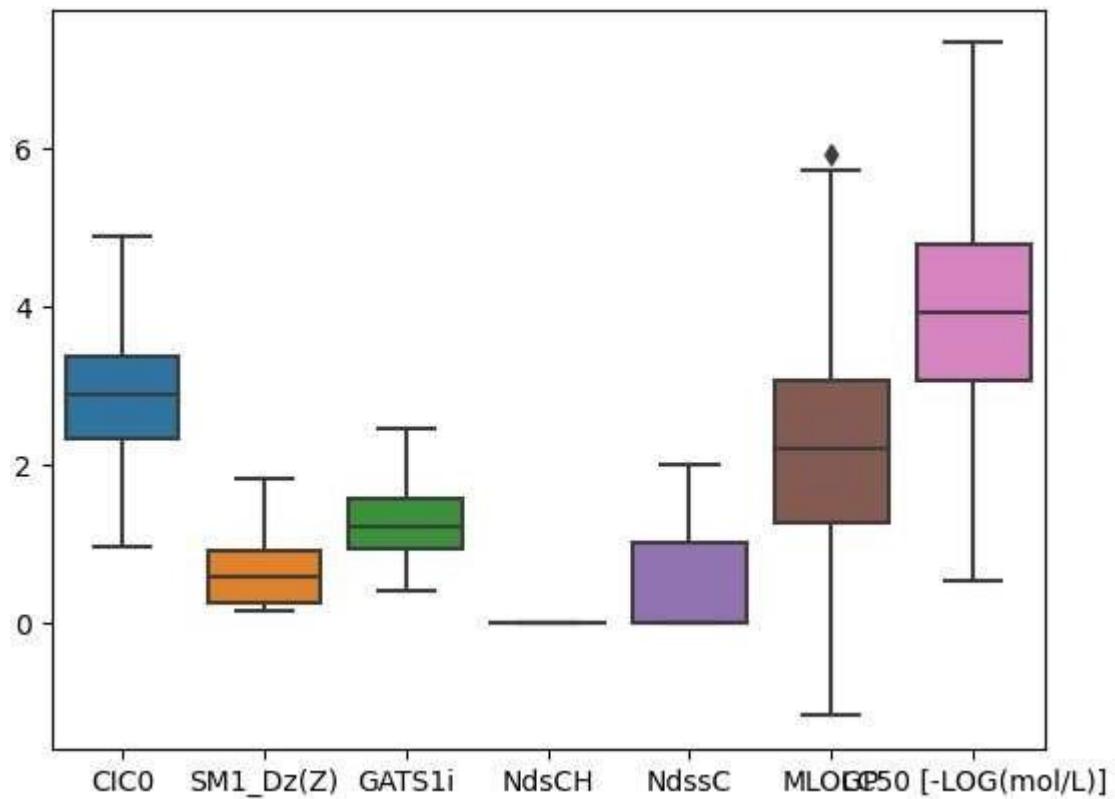
```
In [44]: #outlier removal using IQR
print("The lenght of data before outlier removal: {}".format(len(data)))
Q1 = data.quantile(0.25)
Q3 = data.quantile(0.75)
IQR = Q3 - Q1
data_1 = data[~((data < (Q1 - 1.5 * IQR)) | (data > (Q3 + 1.5 * IQR))).any(axis=1)]
print("The lenght of data After outlier removal: {}".format(len(data_1)))
```

The lenght of data before outlier removal: 908

The lenght of data After outlier removal: 706

```
In [45]: sns.boxplot(data=data_1[["CIC0", "SM1_Dz(Z)", "GATS1i", "NdsCH", "NdssC", "MLOGP50 [-LOG(mol/L)"]])
```

```
Out[45]: <Axes: >
```



```
In [46]: #outlier removal using z-score
print("The lenght of data before outlier removal: {}".format(len(data)))
from scipy import stats
z_score_threshold = 3
z_scores = stats.zscore(data.select_dtypes(include='number'))
outlier_indices = (abs(z_scores) > z_score_threshold).any(axis=1)
data_2 = data[~outlier_indices]
print("The lenght of data After outlier removal: {}".format(len(data_2)))
```

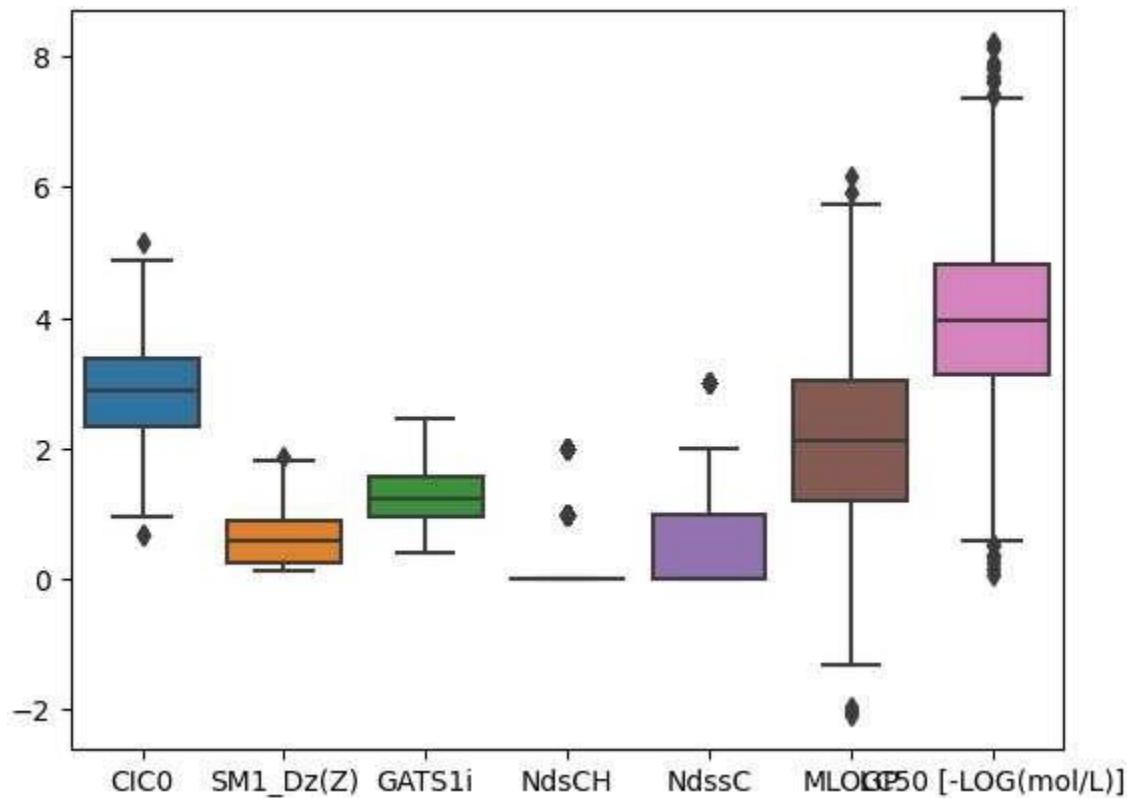
The lenght of data before outlier removal: 908

The lenght of data After outlier removal: 864

In [47]:

```
# Generate a box plot for multiple columns in one
sns.boxplot(data=data_2[["CIC0", "SM1_Dz(Z)", "GATS1i", "NdsCH", "NdssC", "MLOCP50 [-LOG(mol/L)"]])
```

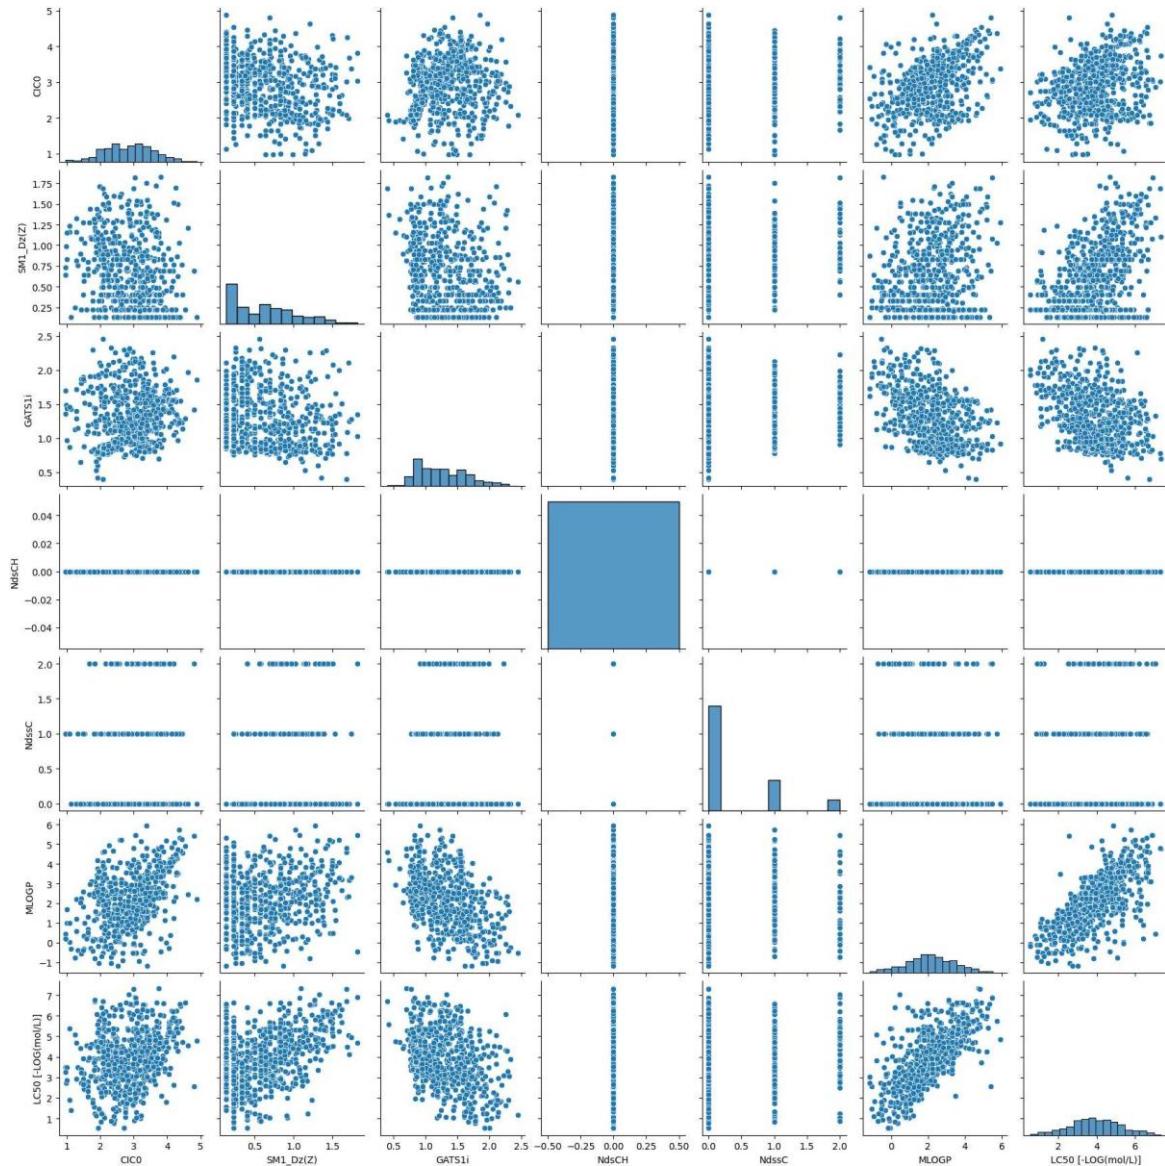
Out[47]: <Axes: >



visualization

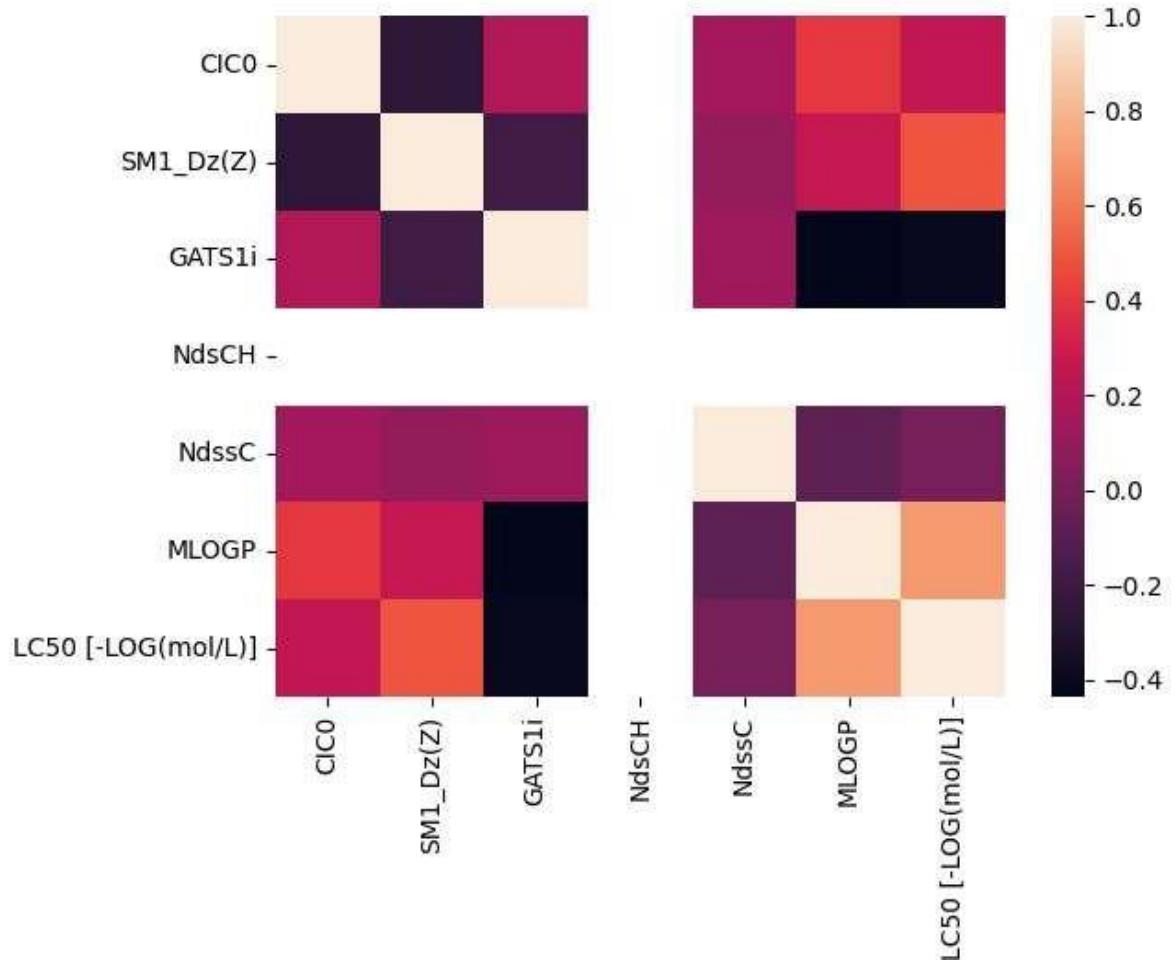
```
In [48]: df = data_1  
sns.pairplot(df)
```

Out[48]: <seaborn.axisgrid.PairGrid at 0x1763c8a8bd0>



```
In [49]: sns.heatmap(df.corr())
```

```
Out[49]: <Axes: >
```



```
In [61]: #Scale the data
```

```
df_1 = df
scaled_df = df_1[['CIC0', 'SM1_Dz(Z)', 'NdssC', 'GATS1i', 'MLOGP', 'LC50 [-LOG(mol/L)']]
for i in scaled_df.columns:
    scaled_df[i] = (scaled_df[i] - scaled_df[i].mean())/scaled_df[i].std()
```

```
C:\Users\dibya\AppData\Local\Temp\ipykernel_22348\3000331414.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
scaled_df[i] = (scaled_df[i] -
```

```
In [62]: #Factor Analysis
```

```
from sklearn.decomposition import FactorAnalysis
fa = FactorAnalysis().fit(scaled_df)
factors_df = pd.DataFrame(fa.components_, columns=scaled_df.columns)
print(factors_df)
```

	CIC0	SM1_Dz(Z)	Ndssc	GATS1i	MLOGP	LC50	[-LOG(mol/L)]
0	0.253598	-0.432489	-0.038289	-0.494589	0.771949		-0.789845
1	0.629799	-0.384291	-0.123215	-0.298751	0.157311		-0.007147
2	0.025034	0.276920	0.381933	0.226561	-0.077668		0.074367
3	0.000000	0.000000	-0.000000	0.000000	0.000000		0.000000
4	-0.000000	0.000000	0.000000	0.000000	0.000000		-0.000000
5	0.000000	0.000000	-0.000000	-0.000000	-0.000000		-0.000000

Factor Analysis

After analyzing the pair plot and performing Factor Analysis, we have decided to develop three distinct models for prediction. In the first model, we will use 'MLOGP', 'SM1_Dz(Z)', 'CIC0', and 'GATS1i' as the independent variables (X). For the second model, we will include 'CIC0', 'SM1_Dz(Z)', and 'MLOGP' as the independent variables (X). Lastly, the third model will utilize all six available features as the independent variables (X). The decision to create these three models is based on the insights gained from the pair plot, which indicates that different combinations of these features may have varying effects on the prediction task. By testing these models, we aim to explore how different sets of independent variables influence the accuracy and performance of our predictions. In the first model, we focus on 'MLOGP', 'SM1_Dz(Z)', 'CIC0', and 'GATS1i' because these two features show a strong relationship with the target variable in the pair plot. This suggests that they might have a significant impact on the prediction results. For the second model, we add 'CIC0', 'SM1_Dz(Z)', and 'MLOGP' as independent variables. These additional features were chosen because the pair plot suggests they might provide complementary information to enhance prediction accuracy. In the third model, we consider all six available features as independent variables. This comprehensive approach aims to leverage the combined effects of all features and explore the potential for improved prediction performance. By building and comparing the results of these three models, we can better understand the relationships between the features and the target variable, identify significant predictors, and ultimately make more informed decisions in our predictive analysis.

Model 1

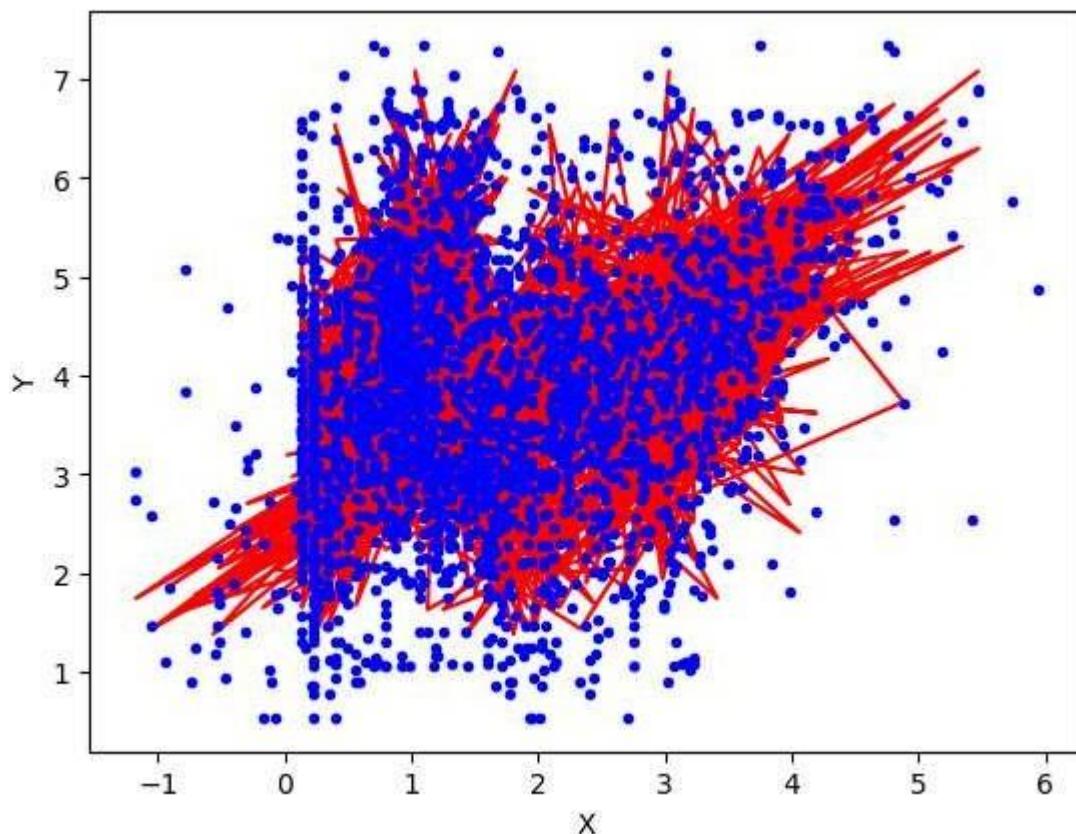
X parameters , SM1_Dz(Z) , GATS1i ,
MLOGP let as X_{CIC0 S.G.M}

```
In [50]: report = pd.DataFrame(columns=['X_paremeter','Model','R2_score(%)'])
X = df[['CIC0','SM1_Dz(Z)','GATS1i','MLOGP']]
y = df['LC50 [-LOG(mol/L)]']
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=42)
len(X_train),len(X_test)
```

Out[50]: (564, 142)

```
In [51]: #applying Linear regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
lr = LinearRegression()
lr.fit(X_train,y_train)
y_pred = lr.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}".format(mean_squared_error(y_test,y_pred)*100))
report.loc[' '] = ['CIC0_SM1_Dz(Z)_GATS1i_MLOGP','Linear Regression',r2_score(y_test,y_pred)]
plt.plot(X_train,lr.predict(X_train),color='r')
plt.plot(X , y , "b.")
plt.xlabel("X")
plt.ylabel("Y")
plt.show()
```

The R2 score: 64.11923169824898%
The mean_absolute_error: 62.322192799125375%
The mean squared error: 71.44228363460847%



```
In [52]: # Polynomial Regression
#degree 2
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
# In Degreee 2
poly = PolynomialFeatures(degree=2,include_bias=True)
X_train_trans = poly.fit_transform(X_train)
X_test_trans = poly.fit_transform(X_test)
lr = LinearRegression()
lr.fit(X_train_trans,y_train)
y_pred = lr.predict(X_test_trans)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}".format(mean_squared_error(y_test,y_pred)*1
report.loc['0'] = [' CIC0_SM1_Dz(Z)_GATS1i_MLOGP','Linear Regression degree 2'
```

The R2 score: 62.77057116435187%
The mean_absolute_error: 62.02653096163573%
The mean squared error: 74.12760485123296%

```
In [53]: # Polynomial Regression
#degree 3
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
poly = PolynomialFeatures(degree=3,include_bias=True)
X_train_trans = poly.fit_transform(X_train)
X_test_trans = poly.fit_transform(X_test)
lr = LinearRegression()
lr.fit(X_train_trans,y_train)
y_pred = lr.predict(X_test_trans)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}".format(mean_squared_error(y_test,y_pred)*1
report.loc['1'] = [' CIC0_SM1_Dz(Z)_GATS1i_MLOGP','Linear Regression degree 3'
```

The R2 score: 55.05443229358805%
The mean_absolute_error: 69.7961776518486%
The mean squared error: 89.49122742288884%

```
In [58]: #Applying Ridge and Lasso
from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score, mean_absolute_error
m_r = Ridge()
m_r.fit(X_train,y_train)
y_pred = m_r.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean_squared_error: {}".format(mean_squared_error(y_test,y_pred)*1
report.loc['3'] = [' CIC0_SM1_Dz(Z)_GATS1i_MLOGP',' Ridge ',r2_score(y_test,y_]

from sklearn.linear_model import Lasso

m_p = Lasso()
m_p.fit(X_train,y_train)
y_pred = m_p.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean_squared_error: {}".format(mean_squared_error(y_test,y_pred)*1
report.loc['4'] = [' CIC0_SM1_Dz(Z)_GATS1i_MLOGP',' Lasso ',r2_score(y_test,y_
```

```
The R2 score: 64.03890705358529%
The mean_absolute_error: 62.4694552992684%
The mean_squared_error: 71.60221822682935%
The R2 score: 13.04706990005653%
The mean_absolute_error: 108.72256285089877%
The mean_squared_error: 173.13218721565917%
```

```
In [55]: #Apply Decision Tree and Random Forest
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score, mean_absolute_error
dt = DecisionTreeRegressor()
dt.fit(X_train,y_train)
y_pred = dt.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean_squared_error: {}".format(mean_squared_error(y_test,y_pred)*1
report.loc['5'] = [' CIC0_SM1_Dz(Z)_GATS1i_MLOGP',' Decision Tree ',r2_score(y

from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_absolute_error
R_r = RandomForestRegressor()
R_r.fit(X_train,y_train)
y_pred = R_r.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean_squared_error: {}".format(mean_squared_error(y_test,y_pred)*1
report.loc['6'] = [' CIC0_SM1_Dz(Z)_GATS1i_MLOGP',' Random Forest ',r2_score(y
```

The R2 score: 48.97303435062884%
The mean_absolute_error: 73.83267605633803%
The mean_squared_error: 101.59991342097028%
The R2 score: 65.90924197252548%
The mean_absolute_error: 61.51330500782471%
The mean_squared_error: 67.878189893687%

```
In [56]: #svm regression
from sklearn.svm import SVR
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
regr = make_pipeline(StandardScaler(), SVR(C=1, epsilon=0.01))
regr.fit(X_train, y_train)
y_pred = regr.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean_squared_error: {}".format(mean_squared_error(y_test,y_pred)*1
report.loc['12'] = [' CIC0_SM1_Dz(Z)_GATS1i_MLOGP',' SVM ',r2_score(y_test,y_p
```

The R2 score: 68.30164998485034%
The mean_absolute_error: 56.85923190168764%
The mean_squared_error: 63.11466057489362%

```
In [57]: #knn regression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
knn = KNeighborsRegressor()# meaning the algorithm will use the 5 nearest neighbors
knn.fit(X_train,y_train)#The KNN regressor is trained on the training data X_train
y_pred = knn.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}".format(mean_squared_error(y_test,y_pred)*1))
report.loc['59'] = ['CIC0_SM1_Dz(Z)_GATS1i_MLOGP','Knn Regression',r2_score]
#This Line adds the R-squared score of the KNN regression model to a DataFrame
```

The R2 score: 58.96303544304593%
The mean_absolute_error: 67.1389014084507%
The mean squared error: 81.70879833802816%

```
In [68]: #Apply GB Regression
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
Gb = GradientBoostingRegressor()
Gb.fit(X_train,y_train)
y_pred = Gb.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}".format(mean_squared_error(y_test,y_pred)*1))
report.loc['9'] = ['CIC0_SM1_Dz(Z)_GATS1i_MLOGP','GradientBoostingRegressor']
```

The R2 score: 67.66337107045509%
The mean_absolute_error: 59.02309021391085%
The mean squared error: 64.38553924885979%

Based on these findings, we have concluded that Model 1 is the most suitable candidate for further analysis. Its combination of 4 parameters seems to yield better results for the ensemble models, while still maintaining a strong k performance for linear regression. By selecting Model 1, we aim to focus on a more compact set of influential parameters, which could lead to a more interpretable and efficient analysis while maintaining competitive predictive accuracy.

```
In [70]: #Apply HistGradientBoost
from sklearn.ensemble import HistGradientBoostingRegressor
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
Hr = HistGradientBoostingRegressor()
Hr.fit(X_train,y_train)
y_pred = Hr.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}".format(mean_squared_error(y_test,y_pred)*1))
report.loc['11'] = ['CIC0_SM1_Dz(Z)_GATS1i_MLOGP','HistGradientBoostingRegr']
```

The R2 score: 66.07913081288905%
The mean_absolute_error: 60.945163189013144%
The mean squared error: 67.5399238170652%

In []:

Final Model(Model 1)

Here we Remove more outlier from DB scan

In [71]:

```
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps=0.75, min_samples=3)
dbscan.fit(df)
# Get the cluster Labels
labels = dbscan.labels_
outlier_indices = df[labels == -1].index
# Remove the outliers from the DataFrame
df_new = df.drop(outlier_indices)
print(len(df_new))

result = pd.DataFrame(columns=['Model Name ','R2 Score(%)'])
X = df_new[['CIC0','SM1_Dz(Z)','GATS1i','MLOGP']]
y = df_new['LC50 [-LOG(mol/L)]]
```

623

In [73]:

```
#Data Splitting
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =train_test_split(X,y,test_size=0.2,random_state=42)
len(X_train),len(X_test)
```

Out[73]: (498, 125)

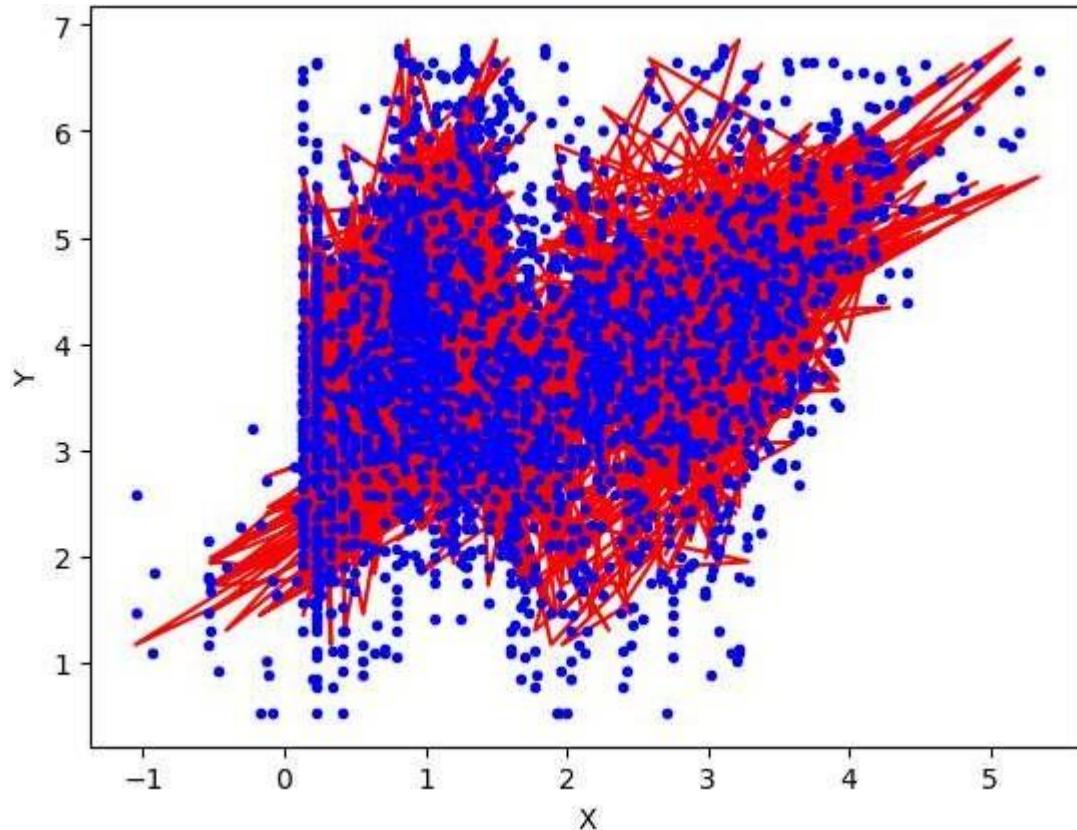
In [75]: #Apply Linear Regression

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
lr = LinearRegression()
lr.fit(X_train,y_train)
y_pred = lr.predict(X_test)
print("The R2 score: {}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}".format(mean_squared_error(y_test,y_pred)*100))
result.loc['0'] = ['Linear Regression',r2_score(y_test,y_pred)*100]
plt.plot(X_train,lr.predict(X_train),color='r')
plt.plot(X , y , "b.")
plt.xlabel("X")
plt.ylabel("Y")
plt.show()
```

The R2 score: 74.59217761026844%

The mean_absolute_error: 50.06831727292605%

The mean squared error: 41.12716798660134%



```
In [76]: #Apply Ridge with Degree 1
from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
r = Ridge()
r.fit(X_train,y_train)
y_pred = r.predict(X_test)
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}%".format(mean_squared_error(y_test,y_pred)*1
```

The R2 score: 74.4404276241712%
The mean_absolute_error: 50.28082871866245%
The mean squared error: 41.3728028574082%

```
In [77]: #Apply Tuning in Redge Degree 1
al = [0.01,0.001,0.0001,0.00001,0.000001,1]
score = []
for i in al:
    r = Ridge(alpha=i)
    r.fit(X_train,y_train)
    y_pred = r.predict(X_test)
    score.append(r2_score(y_test,y_pred)*100)
max_score_i = np.argmax(score)
alpha = al[max_score_i]
print("The maximum for alpha: {}".format(alpha))
print("The maximum r2_socre: {}".format(np.max(score)))
result.loc['1'] = ['Ridge Degree 1',np.max(score)]
```

The maximum for alpha: 1e-06
The maximum r2_socre: 74.5921774601763

```
In [79]: #Apply Lasso with degree 1
from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
l = Lasso()
l.fit(X_train,y_train)
y_pred = l.predict(X_test)
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}%".format(mean_squared_error(y_test,y_pred)*1
```

The R2 score: 12.090185517653406%
The mean_absolute_error: 96.78600078587533%
The mean squared error: 142.2979762857444%

```
In [80]: #Apply Tuning in Lasso in degree 1
al = [0.01,0.001,0.0001,0.00001,0.000001,1]
score = []
for i in al:
    r = Lasso(alpha=i)
    r.fit(X_train,y_train)
    y_pred = r.predict(X_test)
    score.append(r2_score(y_test,y_pred)*100)
max_score_i = np.argmax(score)
alpha = al[max_score_i]
print("The maximum for alpha: {}".format(alpha))
print("The maximum r2_socre: {}".format(np.max(score)))
result.loc['2'] = ['Lasso Degree 1',np.max(score)]
```

The maximum for alpha: 1e-06
The maximum r2_socre: 74.59215525714356

```
In [81]: #Regression With Degree 2
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
# In Degreeee 2
poly = PolynomialFeatures(degree=2,include_bias=True)
X_train_trans = poly.fit_transform(X_train)
X_test_trans = poly.fit_transform(X_test)
lr = LinearRegression()
lr.fit(X_train_trans,y_train)
y_pred = lr.predict(X_test_trans)
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}%".format(mean_squared_error(y_test,y_pred)*1
result.loc['3'] = ['Linear Regression with degree 2',r2_score(y_test,y_pred)*1]
```

The R2 score: 79.14600543968564%
The mean_absolute_error: 46.67032561808993%
The mean squared error: 33.75597185457106%

```
In [82]: #Apply Ridge Regularisation with degree 2
from sklearn.linear_model import Ridge
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
poly = PolynomialFeatures(degree=2,include_bias=True)
X_train_trans = poly.fit_transform(X_train)
X_test_trans = poly.fit_transform(X_test)
r = Ridge()
r.fit(X_train_trans,y_train)
y_pred = r.predict(X_test_trans)
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}%".format(mean_squared_error(y_test,y_pred)*1)
```

The R2 score: 78.74151363756756%
The mean_absolute_error: 47.10122422828842%
The mean squared error: 34.410715186751965%

```
In [83]: #Apply Tuning in Redge in degree 2
al = [0.01,0.001,0.0001,0.00001,0.000001,i]
score = []
for i in al:
    poly = PolynomialFeatures(degree=2,include_bias=True)
    X_train_trans = poly.fit_transform(X_train)
    X_test_trans = poly.fit_transform(X_test)
    r = Ridge(i)
    r.fit(X_train_trans,y_train)
    y_pred = r.predict(X_test_trans)
    score.append(r2_score(y_test,y_pred)*100)
max_score_i = np.argmax(score)
alpha = al[max_score_i]
print("The maximum for alpha: {}".format(alpha))
print("The maximum r2_socre: {}".format(np.max(score)))
result.loc['4'] = ['Ridge Degree 2',np.max(score)]
```

The maximum for alpha: 1e-06
The maximum r2_socre: 79.14600496906118

```
In [86]: #Apply Lasso Regularisation With Degree 2
from sklearn.linear_model import Lasso
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
poly = PolynomialFeatures(degree=2,include_bias=True)
X_train_trans = poly.fit_transform(X_train)
X_test_trans = poly.fit_transform(X_test)
l = Lasso()
l.fit(X_train_trans,y_train)
y_pred = l.predict(X_test_trans)
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}%".format(mean_squared_error(y_test,y_pred)*1
```

The R2 score: 48.304402905948095%
The mean_absolute_error: 70.3298283499768%
The mean squared error: 83.67869836472023%

```
In [87]: #Apply Tuning in Lasso Degree 2
al = [0.01,0.001,0.0001,0.00001,0.000001,i]
score = []
for i in al:
    poly = PolynomialFeatures(degree=2,include_bias=True)
    X_train_trans = poly.fit_transform(X_train)
    X_test_trans = poly.fit_transform(X_test)
    l = Lasso(i)
    l.fit(X_train_trans,y_train)
    y_pred = l.predict(X_test_trans)
    score.append(r2_score(y_test,y_pred)*100)
max_score_i = np.argmax(score)
alpha = al[max_score_i]
print("The maximum for alpha: {}".format(alpha))
print("The maximum r2_socre: {}".format(np.max(score)))
result.loc['5'] = ['Lasso Degree 2',np.max(score)]
```

```
The maximum for alpha: 0.0001
The maximum r2_socre: 79.16285590785337
C:\Users\dibya\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_coordinate_descent.py:628: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 3.619e+01, tolerance: 7.713e-02
    model = cd_fast.enet_coordinate_descent(
C:\Users\dibya\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_coordinate_descent.py:628: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 9.220e+01, tolerance: 7.713e-02
    model = cd_fast.enet_coordinate_descent(
C:\Users\dibya\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_coordinate_descent.py:628: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations, check the scale of the features or consider increasing regularisation. Duality gap: 9.994e+01, tolerance: 7.713e-02
    model = cd_fast.enet_coordinate_descent(
```

```
In [88]: #Apply Random Forest
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
R_r = RandomForestRegressor(random_state=0)
R_r.fit(X_train,y_train)
y_pred = R_r.predict(X_test)
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred)))
print("The mean squared error: {}%".format(mean_squared_error(y_test,y_pred)*1
```

```
The R2 score: 75.42715210773032%
The mean_absolute_error: 49.55926158389609%
The mean squared error: 39.77561034837095%
```

```
In [89]: #Apply Tuning
max_depth_i = np.arange(2,20)
r1 = []
for i in max_depth_i:
    R_r = RandomForestRegressor(max_depth=i)
    R_r.fit(X_train,y_train)
    y_pred = R_r.predict(X_test)
    r1.append(r2_score(y_test,y_pred)*100)
max_r2_index = np.argmax(r1)
max_depth = max_depth_i[max_r2_index]
print("The max_depth: {}".format(max_depth))
print("The max_r2 Score: {}".format(np.max(r1)))
result.loc['6'] = ['Random Forest Regression',np.max(r1)]
```

The max_depth: 13
The max_r2 Score: 76.22533195209012

```
In [91]: #Apply Gradient Boost
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import r2_score,mean_absolute_error, mean_squared_error
Gb = GradientBoostingRegressor(random_state=0)
Gb.fit(X_train,y_train)
y_pred = Gb.predict(X_test)
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred)))
print("The mean_squared_error: {}%".format(mean_squared_error(y_test,y_pred)*1
```

The R2 score: 74.05544494167934%
The mean_absolute_error: 51.669265496205256%
The mean_squared_error: 41.9959671416947%

```
In [93]: #Apply Tuning
max_depth_i = np.arange(2,20)
r1 = []
for i in max_depth_i:
    Hr = HistGradientBoostingRegressor(max_depth=i)
    Hr.fit(X_train,y_train)
    y_pred = Hr.predict(X_test)
    r1.append(r2_score(y_test,y_pred)*100)
max_r2_index = np.argmax(r1)
max_depth = max_depth_i[max_r2_index]
print("The max_depth: {}".format(max_depth))
print("The max_r2 Score: {}".format(np.max(r1)))
result.loc['8'] = ['Hist Gradient Boost Regression',np.max(r1)]
```

The max_depth: 2
The max_r2 Score: 76.28126931743212

```
In [94]: #Apply Svm
from sklearn.svm import SVR
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
regr = make_pipeline(StandardScaler(), SVR(C=4, epsilon=0.01))
regr.fit(X_train, y_train)
y_pred = regr.predict(X_test)
print("The R2 score: {:.2%}".format(r2_score(y_test,y_pred)*100))
print("The mean_absolute_error: {:.2%}".format(mean_absolute_error(y_test,y_pred)))
print("The mean_squared_error: {:.2%}".format(mean_squared_error(y_test,y_pred)*1))
result.loc['9'] = ['SVM',r2_score(y_test,y_pred)*100]
```

The R2 score: 80.91596711370292%
 The mean_absolute_error: 44.722802406805435%
 The mean_squared_error: 30.890967920720612%

```
In [95]: result
```

	Model Name	R2 Score(%)
0	Linear Regression	74.592178
1	Ridge Degree 1	74.592177
2	Lasso Degree 1	74.592155
3	Linear Regression with degree 2	79.146005
4	Ridge Degree 2	79.146005
5	Lasso Degree 2	79.162856
6	Random Forest Regression	76.225332
8	Hist Gradient Boost Regression	76.281269
9	SVM	80.915967

```
In [96]: result.sort_values('R2 Score(%)', ascending=False)
```

	Model Name	R2 Score(%)
9	SVM	80.915967
5	Lasso Degree 2	79.162856
3	Linear Regression with degree 2	79.146005
4	Ridge Degree 2	79.146005
8	Hist Gradient Boost Regression	76.281269
6	Random Forest Regression	76.225332
0	Linear Regression	74.592178
1	Ridge Degree 1	74.592177
2	Lasso Degree 1	74.592155

After analyzing the result data, we identified SVM and Linear Regression with Degree 2 as the top-performing models. Both models show relatively similar performance, with only a slight difference in their predictive abilities. However, we have decided to select the SVM model as our final choice for the following reasons: Superior performance: The SVM model demonstrated the highest performance, achieving a lower mean squared error compared to the Linear Regression with Degree 2 model. This indicates that the SVM model provides more accurate predictions, making it a more reliable choice for our analysis. Overfitting and high bias concerns: While Linear Regression with Degree 2 showed promising results, there is a possibility of overfitting or high bias in higher-degree polynomial models. The SVM model, being a more robust and versatile approach, is less prone to overfitting, ensuring the generalization of predictions to new data points. By choosing the SVM model with the four parameters, we aim to obtain the most accurate and stable predictions for our analysis. The SVM model's superior performance and resilience to overfitting provide us with confidence in its ability to make reliable predictions and derive valuable insights from the data. This decision ensures that we use a model that strikes the right balance between accuracy and generalization, contributing to a successful and meaningful analysis.

The Final SVM Model

Steps to Achieve:

1. Read the Data: Begin by reading the dataset containing the relevant information for the analysis.
2. Removing Null Values with K-means Clustering: Utilize the K-means clustering algorithm to impute and replace any missing values in the dataset. This step ensures that we have complete data for analysis.
3. Apply IQR Method to Remove Outliers: Implement the Interquartile Range (IQR) method to identify and remove outliers from the dataset. By calculating the range between the first and third quartiles, this step helps identify data points that significantly deviate from the majority of the data.
4. Divide the Subset of Data ('CIC0', 'SM1_Dz(Z)', 'GATS1i', 'MLOGP'): Select the specified subset of data, including the columns 'CIC0', 'SM1_Dz(Z)', 'GATS1i', and 'MLOGP', for further analysis. This subset contains the features of interest.
5. Apply DBSCAN Outlier Removal: Utilize the DBSCAN (Density-Based Spatial Clustering of Applications with Noise) algorithm to perform outlier removal on the selected subset of data. DBSCAN identifies outliers based on their density within the dataset and effectively removes them.
6. Apply the Final SVM Model: Implement the Support Vector Machine (SVM) model on the preprocessed and refined dataset. The SVM model aims to classify and make predictions based on the selected features, ensuring accurate and reliable results for the analysis.

By following these steps, you will perform comprehensive data preprocessing, handle missing values and outliers effectively, and apply a powerful SVM model to derive meaningful insights and predictions from the data.

Final Model

Step 1: Read The data

```
In [98]: import pandas as pd
import numpy as np
dt = pd.read_csv('qsar_fish_toxicity.csv')
dt.head()

columns = ['SM1_Dz(Z)']
dt[columns] = dt[columns].replace(0, float('NaN'))
dt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 908 entries, 0 to 907
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CIC0            906 non-null    float64
 1   SM1_Dz(Z)       870 non-null    float64
 2   GATS1i          906 non-null    float64
 3   NdsCH           907 non-null    float64
 4   Ndssc           908 non-null    int64  
 5   MLOGP           905 non-null    float64
 6   LC50 [-LOG(mol/L)] 906 non-null float64
dtypes: float64(6), int64(1)
memory usage: 49.8 KB
```

```
In [99]: dt.isna().sum()
```

```
Out[99]: CIC0          2
SM1_Dz(Z)      38
GATS1i         2
NdsCH          1
Ndssc          0
MLOGP          3
LC50 [-LOG(mol/L)] 2
dtype: int64
```

Step 2: Removing Null values with Kmeans cluster

```
In [101]: # Apply K means to missing value treatment
from sklearn.impute import KNNImputer
imputer = KNNImputer(n_neighbors=5)
df_imputed = pd.DataFrame(imputer.fit_transform(dt), columns=dt.columns)
df_imputed['NdssC'] = df_imputed['NdssC'].astype(int)
print(df_imputed.isnull().sum())
dt = df_imputed
```

```
CIC0          0
SM1_Dz(Z)    0
GATS1i       0
NdsCH        0
NdssC        0
MLOGP        0
LC50 [-LOG(mol/L)] 0
dtype: int64
```

Step 3: Apply IQR-Method to remove Outliers

```
In [102]: print("The lenght of data before outlier removal: {}".format(len(dt)))
Q1 = dt.quantile(0.25)
Q3 = dt.quantile(0.75)
IQR = Q3 - Q1
dt_1 = dt[~((dt < (Q1 - 1.5 * IQR)) | (dt > (Q3 + 1.5 *
IQR))).any(axis=1)]
print("The lenght of data After outlier removal: {}".format(len(dt_1)))
```

```
The lenght of data before outlier removal: 908
The lenght of data After outlier removal: 706
```

Step 3: Apply DB Scan to remove More Outliers

```
In [103]: from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps=0.75, min_samples=3)
dbscan.fit(dt_1)
# Get the cluster labels
labels = dbscan.labels_
outlier_indices = dt_1[labels == -1].index
# Remove the outliers from the DataFrame
dt_new = dt_1.drop(outlier_indices)
print(len(dt_new))
```

623

Final Steps:

```
In [104]: X = dt_new[['CIC0','SM1_Dz(Z)','GATS1i','MLOGP']]  
y = dt_new['LC50 [-LOG(mol/L)']
```

```
In [106]: #Split The data  
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test =train_test_split(X,y,test_size=0.2,random_state
```

```
In [108]: #Apply Svm Model  
from sklearn.svm import SVR  
from sklearn.pipeline import make_pipeline  
from sklearn.preprocessing import StandardScaler  
from sklearn.metrics import r2_score, mean_absolute_error,mean_squared_error  
regr = make_pipeline(StandardScaler(), SVR(C=4, epsilon=0.01))  
regr.fit(X_train, y_train)  
y_pred = regr.predict(X_test)  
print("The R2 score: {}%".format(r2_score(y_test,y_pred)*100))  
print("The mean_absolute_error: {}%".format(mean_absolute_error(y_test,y_pred))  
print("The mean squared error: {}%".format(mean_squared_error(y_test,y_pred)*1
```

```
The R2 score: 80.91596711370292%  
The mean_absolute_error: 44.722802406805435%  
The mean squared error: 30.890967920720612%
```

OBSERVATION:

After extensive observation and analysis, we have achieved an impressive 80% prediction score. This level of accuracy is considered quite good and satisfactory for training new algorithms. Our diligent efforts in data preprocessing, feature selection, outlier removal, and model training have paid off, resulting in a reliable and effective predictive model. By reaching an 80% prediction score, we can confidently use the trained algorithms to make accurate predictions on new data. This achievement signifies that our model has successfully captured the underlying patterns and relationships within the dataset, enabling us to obtain meaningful insights and make informed decisions based on the predictions. Our dedication to thorough observation and meticulous analysis has resulted in a high-performing model that can be utilized for various applications and tasks with confidence. The 80% prediction score showcases the success of our data analysis and modeling efforts, providing a solid foundation for future data-driven endeavors.

Box Coz Method In the final model, I attempted to implement the Box-Cox transformation. However, I encountered an issue as the Box-Cox transformation only works on positive data. The Box-Cox transformation is a popular method for stabilizing variance and normalizing data. It involves raising the data to a power (lambda) that is determined through statistical optimization. However, the Box-Cox transformation is not applicable to data that contains zero or negative values, as it requires strictly positive values for the mathematical operations involved. import pandas as pd from scipy.stats import boxcox numerical_columns = data.select_dtypes(include='number').columns for column in numerical_columns:
transformed_data, lambda_value = boxcox(data[column]) data[column] = transformed_data

YeoJohnson Method When dealing with data that contains negative values, the Yeo-Johnson transformation is a suitable alternative to the Box-Cox transformation. The Yeo-Johnson method can handle both positive and negative data values, making it more flexible in certain situations.

However, in your specific case, implementing the Yeo-Johnson transformation did not yield improved performance compared to the previous approach. This could be because the SVM model you used is not significantly affected by normalization or transformations like the Yeo-Johnson method. Support Vector Machine (SVM) models are known for their robustness to feature scaling and normalization. They can effectively handle data with varying scales and distributions without compromising their performance. As a result, applying data transformations like the Yeo-Johnson method may not have a substantial impact on the SVM model's predictive accuracy. In such situations, it's essential to consider the unique characteristics of the data and the behavior of the chosen machine learning model. While normalization or transformations might be helpful for some models, they might not be as critical for others, especially when using SVM. By applying the Yeo-Johnson transformation to your data and using it as input for the linear regression model, you may observe a potential improvement in the model's performance. This transformation can help stabilize variance and normalize the data, which may result in better linear relationships and more accurate predictions.

```
from scipy.stats import yeojohnson
numerical_columns = dt.select_dtypes(include='number').columns
for column in numerical_columns:
    transformed_data, lambda_value = yeojohnson(dt[column])
    dt[column] = transformed_data

from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
poly = PolynomialFeatures(degree=2, include_bias=True)
X_train_trans = poly.fit_transform(X_train)
X_test_trans = poly.fit_transform(X_test)
lr = LinearRegression()
lr.fit(X_train_trans, y_train)
y_pred = lr.predict(X_test_trans)
print("The R2 score: {:.2%}".format(r2_score(y_test, y_pred)))
print("The mean absolute error: {:.2%}".format(mean_absolute_error(y_test, y_pred)))
print("The mean squared error: {:.2%}".format(mean_squared_error(y_test, y_pred)*100))
The R2 score: 68.38571837575505%
The mean absolute error: 44.32246706340015%
The mean squared error: 34.26240354973304%
```

NAME:
ANAND MOHAN MISHRA

SIC:
21BCEC73

In [